



CSE 512: Distributed and Parallel Database Systems

Lecture 1

Instructor: Mohamed Sarwat

What is a database

- A very large, integrated collection of data.
- Models real-world enterprise.
 - Entities (e.g., students, courses)
 - Relationships (e.g., Mike is taking CSE 412)

Distributed System

A collection of independent computers that appears to its users as a single coherent system.

Distributed System

- Multiple autonomous components
- Components are not shared by all users
- Software runs in concurrent processes on different processors
- Multiple Points of control
- Multiple Points of failure



Why Distribution ?



Multiple Processors

Multiple Memory Systems

Multiple Storage

Why Distribution

Sharing Information and Services

More System Components means:

- More Availability
- More Reliability
- Better Fault Tolerance
- More Scalability

Availability

If a component goes down,
another component is available to
perform the task

Scalability

- Adaption of distributed systems to
 - Accommodate more users
 - Accommodate large volume of data (big data)
 - respond faster (reduce application latency)
- Usually done by adding more and/or faster processors.




Reliability / Fault Tolerance

- Hardware, software and networks fail!
- Distributed systems must maintain availability even at low levels of hardware/software/network reliability.
- Fault tolerance is achieved by
 - Recovery strategy
 - Redundancy (replication)

Transparency

Distributed systems should be perceived by users and application programmers as a whole rather than as a collection of cooperating components.

Distributed DBMS Environment



What is a Distributed Database System?

A distributed database (DDB) is a collection of multiple, *logically interrelated* databases distributed over a *computer network*.

A distributed database management system (D–DBMS) is the software that manages the DDB and provides an access mechanism that makes this distribution *transparent* to the users.

Distributed database system (DDBS) = DDB + D–DBMS



Distributed DBMS Promises

- Transparent management of distributed, fragmented, and replicated data
- Improved reliability/availability through distributed transactions
- Improved performance
- Easier and more economical system expansion

Transparency

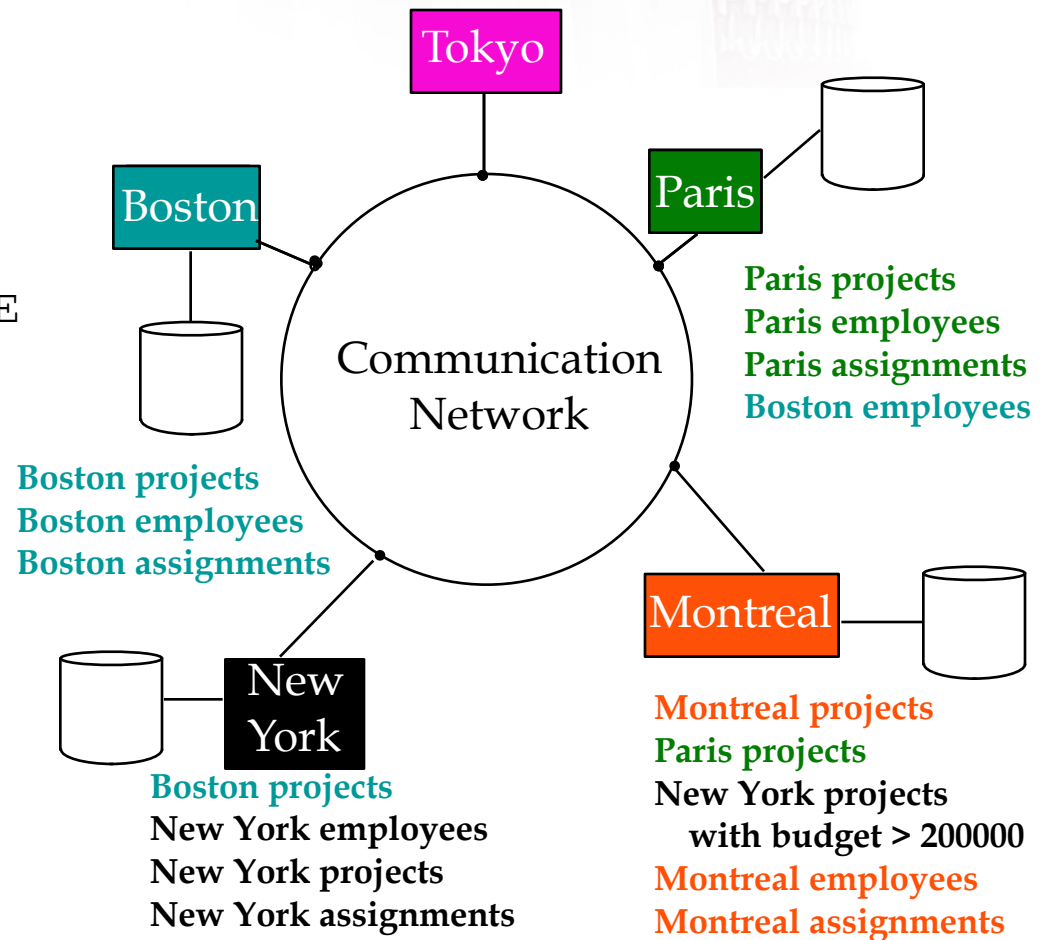
- Transparency is the separation of the higher level semantics of a system from the lower level implementation issues.
- Fundamental issue is to provide **data independence** in the distributed environment



Why Transparency is Key ?

Transparent Access

```
SELECT ENAME, SAL
FROM EMP, ASG, PAY
WHERE DUR > 12
AND EMP.ENO = ASG.ENO
AND PAY.TITLE = EMP.TITLE
```




Storing Data


- Fragmentation
 - Horizontal:
 - Vertical:
- Replication

TID					
t1					
t2					
t3					
t4					

SITE A



SITE B





Fragmentation

Easiest Way to Do Fragmentation ?



Fragmentation Alternatives – Horizontal

$PROJ_1$: projects with budgets less than \$200,000

$PROJ_2$: projects with budgets greater than or equal to \$200,000

PROJ

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal
P2	Database Develop.	135000	New York
P3	CAD/CAM	250000	New York
P4	Maintenance	310000	Paris
P5	CAD/CAM	500000	Boston

$PROJ_1$

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal
P2	Database Develop.	135000	New York

$PROJ_2$

PNO	PNAME	BUDGET	LOC
P3	CAD/CAM	250000	New York
P4	Maintenance	310000	Paris
P5	CAD/CAM	500000	Boston

Fragmentation Alternatives – Vertical

PROJ₁: information about project budgets

PROJ₂: information about project names and locations

PROJ

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal
P2	Database Develop.	135000	New York
P3	CAD/CAM	250000	New York
P4	Maintenance	310000	Paris
P5	CAD/CAM	500000	Boston

PROJ₁

PNO	BUDGET
P1	150000
P2	135000
P3	250000
P4	310000
P5	500000

PROJ₂

PNO	PNAME	LOC
P1	Instrumentation	Montreal
P2	Database Develop.	New York
P3	CAD/CAM	New York
P4	Maintenance	Paris
P5	CAD/CAM	Boston

Which are good fragmentations?

Example:

$$\mathbf{F} = \{ F_1, F_2 \}$$

$$F_1 = \mathbf{\Sigma}_{\text{sal} < 10} E$$

$$F_2 = \mathbf{\Sigma}_{\text{sal} > 20} E$$



Which are good fragmentations?

Example:

$$\mathbf{F} = \{ F_1, F_2 \}$$

$$F_1 = \sigma_{\text{sal} < 10} E$$

$$F_2 = \sigma_{\text{sal} > 20} E$$

➡ Problem: Some tuples lost!

Which are good fragmentations?

Second example:

$$\mathbf{F} = \{ F3, F4 \}$$

$$F3 = \sigma_{\text{sal} < 10 \text{ E}} \quad F4 = \sigma_{\text{sal} > 5 \text{ E}}$$

Which are good fragmentations?

Second example:

$$\mathbf{F} = \{ F3, F4 \}$$

$$F3 = \sigma_{\text{sal} < 10} E \quad F4 = \sigma_{\text{sal} > 5} E$$

➡ Tuples with $5 < \text{sal} < 10$ are duplicated...

⇒ Prefer to deal with replication explicitly

Example: $\mathbf{F} = \{ F_5, F_6, F_7 \}$

$$F_5 = \mathbf{O}_{\text{sal} \leq 5} E \quad F_6 = \mathbf{O}_{5 < \text{sal} < 10} E$$

$$F_7 = \mathbf{O}_{\text{sal} \geq 10} E$$

☞ Then replicate F_6 if convenient
(part of allocation problem)

Desired Properties for Fragmentation

- Completeness
 - Decomposition of relation R into fragments R_1, R_2, \dots, R_n is complete if and only if each data item in R can also be found in some R_i

Desired Properties for Fragmentation

- Reconstruction
 - If relation R is decomposed into fragments R_1, R_2, \dots, R_n , then there should exist some relational operator ∇ such that

$$R = \nabla_{1 \leq i \leq n} R_i$$



Desired Properties for Fragmentation

- Disjointness
 - If relation R is decomposed into fragments R_1, R_2, \dots, R_n , and data item d_i is in R_j , then d_i should not be in any other fragment R_k ($k \neq j$).

Activity

PROJ

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal
P2	Database Develop.	135000	New York
P3	CAD/CAM	250000	New York
P4	Maintenance	310000	Paris
P5	CAD/CAM	500000	Boston

Give an example of a fragmentation that does not satisfy completeness

Give an example of a fragmentation that does not satisfy reconstruction

Give an example of a fragmentation that does not satisfy disjointness



Why Replication ?

When does it become challenging?

Why Replication?

- Gives increased availability.
- Faster query evaluation.
- Updates are Challenging



Potentially Improved Performance

- Proximity of data to its points of use
 - Requires some support for fragmentation and replication